

Data Diversity: An Approach to Software Fault Tolerance

PAUL E. AMMANN AND JOHN C. KNIGHT

Abstract—Crucial computer applications require extremely reliable software. For a typical system, current proof techniques and testing methods cannot guarantee the absence of software faults, but careful use of redundancy may allow the system to tolerate them. Existing methods to provide fault tolerance at execution time rely on redundant software written to the same specifications. Such techniques use *design diversity* to tolerate residual faults.

Diversity in the data space can also provide fault tolerance. This is because program faults often cause failure only under certain special case conditions, and for some applications a program may express its input and internal state in a large number of logically equivalent ways. These observations suggest obtaining a related set of points in the data space, executing the same software on these points, and then employing a decision algorithm to determine system output. Such techniques use *data diversity* to tolerate residual faults.

Index Terms—Data diversity, design diversity, *N*-copy programming, *N*-version programming, recovery blocks, retry blocks, software faults, software fault tolerance.

I. INTRODUCTION

RESEARCHERS have proposed various methods for building fault-tolerant software in an effort to provide substantial improvements in the reliability of software for crucial applications. At execution time, the fault-tolerant structure attempts to cope with the effect of those faults that survive the development process. The two best-known methods of building fault-tolerant software are *N*-version programming [3] and recovery blocks [11]. To tolerate faults, both of these techniques rely on *design diversity*, the availability of multiple implementations of a specification. Software engineers assume that the different implementations use different designs and thereby, it is hoped, contain different faults. The fact that diversity in the design space may provide fault tolerance suggests that diversity in the data space might also. This paper considers *data diversity* [1], [2], a fault-tolerant strategy that complements design diversity.

N-version programming requires the separate, independent preparation of multiple (i.e., "*N*") versions of a program for some application. These versions execute in parallel in the application environment; each receives identical inputs, and each produces its version of the required outputs. A voter collects the outputs, that should, in principle, all be the same.

Manuscript received June 8, 1987; revised November 23, 1987. This work was supported in part by NASA under Grant NAG-1-605.

The authors were with the Department of Computer Science, University of Virginia, Charlottesville, VA 22903. They are now with The Software Productivity Consortium, Reston, VA 22091.

IEEE Log Number 8719329.

If the outputs disagree, the system uses the results of the majority, provided there is one.

The recovery block structure submits the results of an algorithm to an acceptance test. If the results fail the test, the system restores the state of the machine that existed just prior to execution of the algorithm and executes an alternate algorithm. The system repeats this process until it exhausts the set of alternates or produces a satisfactory output.

It is well known that software often fails for special cases in the data space; for example, see [12]. In practice, a program may survive extensive testing, work for many cases, and then fail on a special case. The special case may take the form of what seems to be an obscure set of values in the data. Testing frequently fails to reveal faults associated with special cases precisely because the test harness does not generate the exact circumstances required. A test data set whose values are merely close to the values that cause the program to fail does not uncover the fault.

These observations suggest that if software fails under a particular set of execution conditions, a minor perturbation of those execution conditions might allow the software to work. Other researchers have exploited this property in specific instances. Gray observed that certain faults that caused failure in an asynchronous commercial system did not always cause failure if the same inputs were submitted to a second execution [6]. The system succeeded on the second execution due to a chance reordering of the asynchronous events. Gray introduced the term "Heisenbugs" to describe these faults and their apparent nondeterministic manifestations.

Shepherd, Martin, and Morris have proposed "temporal separation" of the input data to a dual version system [8], [9]. The versions use data from adjacent real-time frames rather than the same frame. Since the versions read sensor data at different times, the data tend to differ. The system uses old outputs to evaluate the validity of current outputs. It is hoped that the use of time-skewed data will allow the system to select the correct version in the event of failure.

Each of these approaches attempts to avoid faults by operating software with altered execution conditions. Each approach relies upon circumstance to change the conditions. However, execution conditions can be changed deliberately. For example, concurrent systems need not rely on a chance reordering of events. If reordering events might allow a second execution to succeed, then the system should enforce a reordering. Changing the processor dispatching algorithm after state restoration forces a different execution sequence. Similarly, skewing the inputs to the versions in an *N*-version system does not require the passage of time. Inputs can be

manipulated algorithmically. Many real-valued quantities have tolerances set by their specifications, and all values within those tolerances are logically equivalent.

Data diversity is an orthogonal approach to design diversity and a generalization of the work cited above. A diverse-data system produces a set of related data points and executes the same software on each of these points. A decision algorithm then determines system output.

This paper describes data diversity as an approach to fault-tolerant software and presents the results of a pilot study. Section II discusses the regions of the input space that cause failure for certain experimental programs. Section III examines *data reexpression*, the way in which alternate input data sets can be obtained. Section IV describes the *retry block*, the data-diverse equivalent of the recovery block, and presents a model of the retry block together with some empirical results. Section V describes *N-copy programming*, the data-diverse equivalent of *N-version programming*, presents a simple model, and gives some empirical results. Section VI contains conclusions.

II. FAILURE REGIONS

The input data for most programs come from hyperspaces of very high dimension. For example, a program may read and process a set of 20 floating-point numbers, and so its input space has 20 dimensions. In many cases, the number of dimensions in the space varies dynamically because the amount of data that a program processes varies for different executions.

The *failure domain* of a program is the set of input points that cause program failure [5]. We call a failure domain along with its geometry a *failure region*. A failure region describes the distribution of points in the failure domain and determines the effectiveness of data diversity. The fault tolerance of a system employing data diversity depends upon the ability of the reexpression algorithm to produce data points that lie outside of a failure region, given an initial data point that lies within a failure region. The program executes correctly on reexpressed data points only if they lie outside a failure region. If the failure region has a small cross section in some dimension(s), then reexpression should have a high probability of translating the data point out of the failure region.

Knowledge of the geometry of failure regions gives insight into the possible performance of data diversity. We have obtained two-dimensional cross sections of several failure regions for faults in programs used in a previous experiment [7]. The programs are simple hypothetical antimissile missile launch decision programs that determine their output based on relationships among sets of points from an imaginary radar track. For example, one test determines whether certain sets of radar points fit into a circle of given radius.

Fig. 1 shows cross sections obtained by varying two inputs across a uniform grid while all other inputs remained fixed. The cross sections are from two separate faults.¹ The solid lines show where the correct output of the program changes, and the small dots show grid points where the faulty program produces the wrong output.

¹ The specific faults are 6.2 and 6.3 [4].

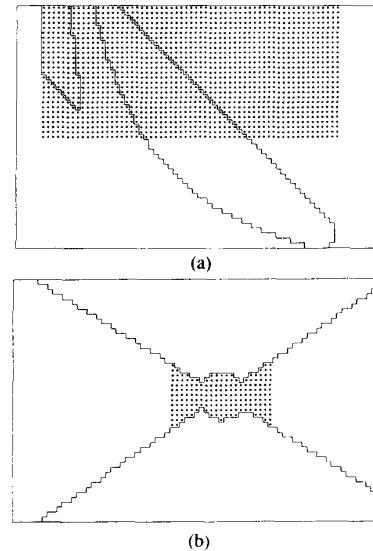


Fig. 1. Two-dimensional cross sections of two failure regions.

The (x, y) coordinates of a set of points in the radar track form the input space. Since each (x, y) pair supplies two dimensions, the input space has twice as many dimensions as radar points. The radar points are distributed so that the x and y coordinates each span the real range $-40 \cdots 40$,² and all radar points for the original experiment were rounded to the nearest 0.1.

The space from which cross section (a) was taken has 30 dimensions corresponding to 15 radar points. Cross section (a) was obtained by varying coordinates x_1 and x_9 with a grid point separation of 0.2. The space from which cross section (b) was taken has 18 dimensions corresponding to nine radar points. Cross section (b) was obtained by varying coordinates x_6 and y_6 with a grid point separation of 0.000001. For both, all other inputs were held fixed. The area of cross section (a) is 4×10^{10} times larger than the area of cross section (b).

The cross sections shown are typical for these programs. This small sample illustrates several important points. First, at the resolution used in scanning, these particular failure regions are locally continuous. Second, since the failure regions vary greatly in size, exiting them varies greatly in difficulty. Finally, failure regions tend to be associated with transitions in the output [1].

III. DATA REEXPRESSION

At its simplest, *data reexpression* is the generation of logically equivalent data sets. Fig. 2 illustrates this basic form of data reexpression. An input x given directly to a program P produces an output $P(x)$. Alternatively, a reexpression algorithm R transforms the original input x to produce a new input y , where $y = R(x)$. The input y may contain exactly the same information as the input x , but in a different form, or y may approximate the information in input x . The program P operates on the reexpressed input y to produce $P(y)$. P and R determine the relationship between $P(x)$ and $P(y)$. Data

² The distribution is not uniform and is defined in [10].




Fig. 2. Data reexpression.

diversity can tolerate faults when $P(y)$ is a useful output but $P(x)$ is not.

Since it is the outputs that are ultimately important to any given application, the requirements for a reexpression algorithm can be derived from characteristics of the outputs. For a specific input x , Fig. 3 shows the sets of input points y , that are of interest given a reexpression algorithm of the type shown in Fig. 2.


One set in Fig. 3, called I , the *identical output set*, is defined $I = \{y \mid \text{Correct}(x, P_{\text{correct}}(y))\}$ and contains all inputs y for which the correct output is the same as the correct output given x as an input. The predicate $\text{Correct}(in, out)$ is true iff out is the output demanded by the specification for input in . The symbol P_{correct} represents a correctly implemented specification; $P_{\text{correct}}(x)$ is the correct output for input x . The identical output set is a desirable set from which to select reexpressed inputs y , since the use of data reexpression in such cases is transparent outside the program. Implementing an error detection mechanism is simplified when all reexpressed data points are members of the identical output set.

It is possible that the single output produced by P_{correct} for a specific input x is not the only acceptable output for the application. Different acceptable outputs may arise at boundaries in the output space. For example, if an input x is a real variable of limited precision, then different outputs resulting from inputs sufficiently near a boundary are indistinguishable from the viewpoint of the application.

The identical output set I is a subset of the *valid output set* V , $V = \{y \mid \text{Valid}(x, P_{\text{correct}}(y))\}$. The predicate $\text{Valid}(in, out)$ is true iff out is a valid or acceptable output for input in . V is the set of all inputs y for which a correct program produces an acceptable output given x as the actual input. Although it is easier to produce values in V than in I , error detection with members of V is more difficult due to the problem of voting with multiple correct outputs.

The third set F , called the *failure set*, is $F = \{y \mid \text{not Valid}(y, P(y))\}$, and represents all inputs y for which the program fails to produce a correct or acceptable output. Although elements of F are, by definition, not enumerated, the effectiveness of data diversity is determined by the proportion of reexpressed points that lie in F .

Fig. 3 identifies two classes of reexpression algorithms. The first class produces elements in I (up to numerical error); these algorithms are termed *exact*. The second class yields elements in V ; these algorithms are termed *approximate*. Although exact algorithms are desirable from the viewpoint of error detection, approximate algorithms may be easier to produce and may have a better chance of escaping a failure region. Exact reexpression algorithms may have the defect of preserving precisely those aspects of the data that cause failure.

Fig. 3. Sets in the output space for given input x .

As an example of an exact reexpression algorithm, suppose that a program processes Cartesian input points and that only the points' relative positions are relevant to the application. A legitimate reexpression algorithm could translate the coordinate system to a new origin or rotate the coordinate system about an arbitrary point.

Any mapping of a program's data that preserves the information content of the data is a valid reexpression algorithm. A simple approximate data reexpression algorithm for a real variable might alter its value by a small percentage. The allowable percentage by which the variable's value can be altered would be determined by the specification, possibly as a result of a known sensitivity in a particular hardware sensor.

Up to now, we have considered reexpression algorithms that depend on numeric manipulation. Data can take other forms, however, and data diversity can be applied successfully to other applications. Consider a compiler with a conventional multipass organization. Although the initial representation of a source program is a character string, the program may be represented in many ways during compilation, for example, as a tree. There are many transformations that can be applied to a tree that preserve semantics. A compiler employing data diversity for its later passes could be constructed by executing these later passes on several different copies of the tree obtained by semantics preserving transformations.

Similarly, the order of storage allocation in an activation record is usually determined by the programmer's order of declaration. In practice, this order need not be preserved, and a set of semantically equivalent internal representations of a program can be obtained by shuffling the order of variables in activation records.

Combining tree transformations, data storage reordering, and code storage reordering, i.e., generation of code for subprograms in an arbitrary order, provides considerable diversity in the data processed by large fractions of a conventional compiler. These approaches to reexpression are exact in that, although the code generated by the compiler may be different and therefore not amenable to any simple selection algorithm, the effects of these programs should be identical, and so simple voting can be used for selection during execution of the programs generated by the compiler.

The structure shown in Fig. 4 allows the reexpression algorithm to produce more diverse inputs than is possible with the structure shown in Fig. 2. Some form of correction, A , is performed on the output after the program is run to undo the distortion caused by reexpression. If the distortion induced by




Fig. 4. Reexpression with postexecution adjustment.

R can be removed after execution, the approach permits major changes to the inputs and allows copies of the program to operate in widely separated regions of the input space. As an example, consider a program which computes intersections of line segments. Reexpression might alter the representation of the input by multiplying by a nonsingular matrix. The distortion could be recovered by multiplying the program output by the inverse of the matrix.

Another approach to reexpression is shown in Fig. 5. An input may be decomposed into a related set of inputs. The program is run on each part of the input, and then the results are recombined. Again, as in Fig. 4, Fig. 5 allows for both exact and approximate reexpression algorithms.

As an example of an exact reexpression algorithm of the type shown in Fig. 5, consider a data-diverse computation of the sine function. Assume we have an implementation of the sine function that is known to work over a large percentage of its inputs; the failure probability for the sine function on a randomly chosen input x is p , where $p \ll 1$. To compute $\sin(x)$, we use the two trigonometric identities

$$\sin(a+b) = \sin(a)\cos(b) + \cos(a)\sin(b)$$

$$\cos(a) = \sin(\pi/2 - a)$$

to rewrite

$$\sin(x) = \sin(a)\sin(\pi/2 - b) + \sin(\pi/2 - a)\sin(b)$$

in which a and b are any two real numbers such that $a + b = x$. Suppose that $\sin(x)$ is computed using three independent decompositions for x obtained by using three different values each for a and b , and that a simple majority voter selects the output. Using the worst case assumption that all incorrect answers appear identical to the voter, a conservative estimate for the probability of computing an incorrect value for $\sin(x)$ can be shown to be on the order of $48p^2$ [1].

In general, a reexpression algorithm must be tailored to the application at hand. Producing a reexpression algorithm requires a careful analysis of the type and magnitude of reexpression appropriate for each candidate datum. Simpler reexpression algorithms are more desirable than complex ones since they are less likely to contain design flaws.

IV. RETRY BLOCKS

A *retry block* is a modification of the recovery block structure that uses data diversity instead of design diversity. Fig. 6 shows the semantics of a retry block. Rather than the multiple alternate algorithms used in a recovery block, a retry block uses only one algorithm. A retry block's acceptance test has the same form and purpose as a recovery block's acceptance test. A retry block executes the single algorithm




Fig. 5. Reexpression via decomposition and recombination.




Fig. 6. Retry block.

normally and evaluates the acceptance test. If the acceptance test passes, the retry block is complete. If the acceptance test fails, the algorithm executes again after the data have been reexpressed. The system repeats this process until it violates a deadline or produces a satisfactory output.

A. Simple Model for a Retry Block

Fig. 7 shows the model we have used to predict the success of a retry block assuming a perfect acceptance test. On a single execution under operational conditions, the program used in the construction of the retry block has a probability of failure, p . The random variable Q gives the probability that a reexpressed data point causes failure given that the initial point caused failure. Q is a random variable because its value depends upon the geometry of the failure region in which the original data point lies, the location of the data point within that region, and the algorithm used to reexpress the data. Since Q is a probability, it assumes values in the range $0 \leq Q \leq 1$. We denote the distribution function for Q as $F_Q(q)$ where $F_Q(q) = \text{prob}(Q \leq q)$. The corresponding density function for Q is $f_Q(q)$ where $f_Q(q) = (d/dq)F_Q(q)$. We assume that Q is continuously distributed. The extension to discretely distributed Q is straightforward.

The probability that the retry block fails when using n retries, denoted $\text{prob}(\text{fail})$, is $\text{prob}(\text{fail}) = pQ^n$. The probability that a retry block will succeed in n or fewer retries is one minus this probability. $\text{prob}(\text{fail})$ is a random variable since it is a function of Q . Its expected value is

$$E[\text{prob}(\text{fail})] = p \int_0^1 q^n f_Q(q) dq.$$

Since the integral in the expression for $E[\text{prob}(\text{fail})]$ is multiplied by p , the failure probability of the program, the integral describes how the performance of a retry block improves upon the performance of a program. Thus, the quantity $\int_0^1 q^n f_Q(q) dq$ is the average factor by which the use of a retry block reduces the probability of failure. We next estimate $F_Q(q)$ for some of the programs discussed in Section II and approximate the integral in $E[\text{prob}(\text{fail})]$ to predict the




Fig. 7. Retry block model assuming a perfect acceptance test.

reductions in failure probability for retry blocks constructed from those programs.

B. Two Forms Of Acceptance Test

In this experiment we have collected data for the performance of the retry block under two different implementations of an acceptance test with an approximate reexpression algorithm. First we consider an acceptance test that can determine whether the program has produced $P_{\text{correct}}(x)$ given x as input. If the acceptance test rejects the program's output for input x and x is reexpressed as y , then the acceptance test demands $P_{\text{correct}}(y)$ as the only acceptable output. This acceptance test rejects all outputs for which the program does not satisfy the specification. We call this type of acceptance test a *specification* acceptance test.

The second form of acceptance test that we consider accepts any output that is valid anywhere in the reexpression region of the input space. The reexpression region is that part of the input space from which reexpressed inputs can be selected. The acceptance test determines the validity of the output by examining the output of the P_{correct} for input points in the reexpression region of input x . We call this type of acceptance test a *valid output* acceptance test.

The first set of empirical results for the retry block was obtained using a specification acceptance test; the second set was obtained using a valid output acceptance test.

C. Empirical Results for a Retry Block—Specification Acceptance Test

We have obtained empirical evidence of the expected performance of data diversity on some of the known faults in the Launch Interceptor Programs produced for the Knight and Leveson experiment [7]. As noted in Section II, one of the inputs to the programs is a list of (x, y) pairs representing radar tracks. To employ data diversity in this application, we assume that data obtained from the radar are of limited precision. The data reexpression algorithm we used moved each (x, y) point to a random location on the circumference of a circle centered at (x, y) and of some small, fixed radius. The reexpression algorithm is approximate according to the definitions given in Section III. Fig. 8 shows how this algorithm reexpresses a set of three radar points. Many other reexpression algorithms are possible.

The experiment measured the performance of data diversity in the form of the retry block on the Launch Interceptor Programs, and how this performance was affected by two parameters. The first parameter studied was the radius of displacement used in the data reexpression algorithm. The second parameter was the effect of the reexpression algorithm on different faults.






Fig. 8. Reexpression of three radar points.

Fig. 9. Fault 9.1 sample $F_Q(q)$ for three displacement values.Fig. 10. Sample $F_Q(q)$ for three different faults at 0.01 displacement.

Figs. 9 and 10 show the effects of these two parameters on estimated $F_Q(q)$ functions. Each distribution function in these figures corresponds to a particular displacement value and fault. A given point $(q, F_Q(q))$ is the observed probability $F_Q(q)$ that a program executing on a reexpressed data point will have at most a probability of failure, q . Distribution functions that rise rapidly imply better performance for data diversity since they indicate a higher probability that reexpression will arrive at a point outside the failure region. For example, a function containing the point $(0.05, 0.95)$ means that the probability is 0.95 that the reexpressed data point will cause failure with a probability of 0.05 or less. Each data point in Figs. 9 and 10 was obtained by averaging the results of 50 applications of the reexpression algorithm to a randomly selected point from a failure region. Fig. 9 has 47 such data points on each of its three graphs; Fig. 10 has 100, 71, and 100, respectively.

Fig. 9 shows the observed $F_Q(q)$ of a single fault for three values for the radius of displacement. From left to right on the graph, the displacement values are 0.1, 0.01, and 0.001. As would be expected, larger displacement values in the data reexpression algorithm have the effect of making the reexpressed data point less likely to cause failure. These displacements are relatively small compared to the range of values that the radar points could assume.

Fig. 10 shows the observed $F_Q(q)$ for three different faults³

³ From left to right on the graph, the faults are 8.1, 7.1, and 6.2 [4].

Retries	1			2			3		
Displacement	0.001	0.01	0.1	0.001	0.01	0.1	0.001	0.01	0.1
Fault									
6.1	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
6.2	1.00	0.98	0.87	1.00	0.96	0.81	0.99	0.94	0.75
6.3	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
7.1	0.92	0.59	0.26	0.87	0.43	0.11	0.80	0.29	0.03
8.1	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
8.2	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
9.1	0.99	0.90	0.39	0.97	0.83	0.19	0.97	0.74	0.07

Fig. 11. Failure probability multipliers using specification acceptance test.

using a fixed displacement of 0.01 in the reexpression algorithm. These three faults were chosen to show the wide variation from fault to fault on the distribution of the probability that a reexpressed data point will cause failure. The leftmost function rises rapidly, which indicates that data diversity will tolerate the associated fault well at the given displacement value. The rightmost function rises slowly, which indicates that data diversity requires a better reexpression algorithm to tolerate the associated fault.

The table in Fig. 11 shows the performance obtained using retry blocks with a specification acceptance test for various faults under various conditions. Each table entry is an expected value multiplier for the probability of system failure. The table shows results for retry blocks with 1, 2, or 3 retries and displacement values of 0.1, 0.01, or 0.001. An entry in Fig. 11 means that the failure probability associated with a given fault is reduced by the factor shown.⁴ For instance, the value of 0.43 found in the middle of the table means that the failure probability associated with fault 7.1 was reduced by a factor of 0.43 when the displacement in the reexpression algorithm was 0.01 and the number of retries was 2. Similarly, a table entry of 0.00 means that the effects of the associated fault were eliminated, and an entry of 1.00 means that data diversity had no effect.

The model outlined earlier in this section was used to calculate the values shown in Fig. 11. The integral corresponding to the multiplier was approximated using sample $F_Q(q)$, examples of which appear in Figs. 9 and 10.

D. Empirical Results for a Retry Block—Valid Output Acceptance Test

Fig. 12 gives the expected value multipliers for the probability of system failure for retry blocks using the valid output acceptance test for 1, 2, and 3 retries. The entries in Fig. 12 were derived in a similar manner to the entries in Fig. 11. The same data yielded both figures, and only the definition of the acceptance test and the meaning of "failure" were different. In both cases, the definition of failure was used consistently when comparing single version performance to the retry block performance. When evaluating the retry block using the specification acceptance test, outputs were considered to be wrong if they disagreed with the specification. Similarly, when evaluating the retry block using the valid output acceptance test, outputs were considered to be wrong if

Retries	1			2			3		
Displacement	0.001	0.01	0.1	0.001	0.01	0.1	0.001	0.01	0.1
Fault									
6.1	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
6.2	1.00	0.99	0.96	1.00	0.98	0.93	1.00	0.98	0.90
6.3	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
7.1	0.96	0.83	0.26	0.93	0.71	0.12	0.90	0.63	0.07
8.1	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
8.2	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
9.1	0.99	0.93	0.70	0.98	0.89	0.54	0.97	0.86	0.44

Fig. 12. Failure probability multipliers using valid output acceptance test.

they did not match the output demanded by the specification for any input points in the reexpression region.

The performance of the retry block did not depend significantly on the type of acceptance test used. In both cases, four of the faults were tolerated completely, and three of the faults were not well tolerated. Although the more generous definition of failure used for the valid output acceptance test results in a lower failure rate than the more stringent definition of failure required by the specification acceptance test, the performance improvement attributable to the retry block is comparable for the two cases.

V. N-COPY PROGRAMMING


An N -copy system is similar to an N -version system but uses data diversity instead of design diversity. Fig. 13 shows the structure of an N -copy system. N copies of a program execute in parallel, each on a set of data produced by reexpression. The system selects the output to be used by an enhanced voting scheme.

Voting in an N -copy system is not necessarily straightforward. If the reexpression algorithm is exact, so that all copies should generate identical outputs, then a conventional majority vote can be used. However, if the reexpression algorithm is approximate, the copies could produce different but acceptable outputs. This is very likely to occur near boundaries in the output space. In that case, a simple majority may not exist and the selection process is more involved. If a particular output occurs more than once, then it might be selected. If more than one output occurs more than once or no output occurs more than once, then selection might have to involve an arbitrary choice.

A. Simple Model for N -Copy Programming

To determine the performance of data diversity as represented by an N -copy system, we analyze a three-copy system and compare it to a single version. The copies can map their inputs to $R + W$ different outputs. R of the outputs are acceptable to the application, and the probability that a copy produces acceptable output i is r_i for $i = 1 \cdots R$. Similarly, W of the outputs are not acceptable, and the probability that a copy produces unacceptable output i is w_i for $i = 1 \cdots W$. To produce an output, the system generates three reexpressed data sets and executes a copy on each. The system selects the output that occurs most frequently, if there is one. In the case of a tie, the tie is broken by choosing at random. Given the probabilities r_i and w_i , we wish to know with what probability the

⁴ The original failure probabilities for the faults are given in [4], [7].

Fig. 13. *N*-copy programming.

system will select an acceptable output. There are four cases under which the system makes an acceptable choice.

1) The program maps all three reexpressed data points to acceptable outputs. This happens with probability $\sum_{i=1}^R \sum_{j=1}^R \sum_{k=1}^R r_i r_j r_k$.

2) The program maps two reexpressed data points to the same acceptable output, but maps one to an unacceptable output. This happens with probability $3 \sum_{i=1}^R \sum_{j=1}^W r_i^2 w_j$.

3) The program maps two reexpressed data points to different acceptable outputs, and maps one to an unacceptable output. Two thirds of the time, an acceptable output is selected at random. The probability of this event is $2 \sum_{i=1}^R \sum_{j=1}^R \sum_{k=1}^W r_i r_j w_k$, such that $i \neq j$.

4) Finally, the program maps one reexpressed data point to an acceptable output, but maps the other two to different unacceptable outputs. One third of the time, an acceptable output is selected at random. The probability of this event is $\sum_{i=1}^R \sum_{j=1}^W \sum_{k=1}^W r_i w_j w_k$, such that $j \neq k$.

The sum of these four probabilities gives the probability that a three-copy system will yield an acceptable output.

B. Empirical Results For *N*-Copy Programming

Fig. 14 gives the expected value multipliers for the probability of system failure for *N*-copy systems using three and five copies. For both three- and five-copy systems, four of the faults were tolerated completely, and three of the faults were not well tolerated. The *N*-copy systems tolerated each fault with about the same success as the retry blocks.

As with the retry block using the valid output acceptance test, the definition of failure for an *N*-copy system was used consistently when comparing single version performance to the *N*-copy system performance. In both cases, outputs were considered to be wrong only if they did not match the output demanded by the specification for any input points in the reexpression region.

The data shown here measure the performance of an *N*-copy system only on failure points and not on a representative sample of input data. Although we expect it to be small, in this experiment we did not measure the probability that an *N*-copy system would fail when operating on points for which a single version succeeds.

VI. CONCLUSIONS

We have described the general concept of data diversity as a technique for software fault tolerance and have defined the retry block and *N*-copy programming as two possible approaches to its implementation. We have presented the results of a pilot study of data diversity for both structures. Although the overall performance of data diversity varied greatly, we observed a large reduction in failure probability

Copies	3			5		
Displacement	0.001	0.01	0.1	0.001	0.01	0.1
Fault						
6.1	0.00	0.00	0.00	0.00	0.00	0.00
6.2	1.00	0.99	0.97	1.00	1.00	0.97
6.3	0.00	0.00	0.00	0.00	0.00	0.00
7.1	0.97	0.88	0.22	0.98	0.91	0.19
8.1	0.00	0.00	0.00	0.00	0.00	0.00
8.2	0.00	0.00	0.00	0.00	0.00	0.00
9.1	0.99	0.95	0.73	0.99	0.96	0.74

Fig. 14. Multipliers for probability of system failure for three and five copies.

for some of the faults examined in the study. In several cases, data diversity completely eliminated the effects of a fault.

The success of data diversity depends, in part, upon developing a data reexpression algorithm that is acceptable to the application, yet has a high probability of generating data points outside of a program's failure region. Many applications could use simple reexpression algorithms similar to the one employed in this study. For example, sensors typically provide data with relatively little precision and small modifications to those data would not affect the application.

An issue in the implementation of a retry block is the need for a suitable acceptance test. This is a well-known problem for the recovery block, and any techniques developed for the recovery block apply directly to the retry block.

Compared to design diversity, data diversity is relatively easy and inexpensive to implement. Data diversity requires only a single implementation of a specification, although additional costs are incurred in the data reexpression algorithm.

Data diversity is orthogonal to design diversity. The strategies are not mutually exclusive and various combinations are possible. For example, an *N*-version system in which each version was an *N*-copy system could be built for very little additional cost over an *N*-version system. One possible advantage of integrating design and data diversity is that the strategies may remedy different classes of errors. The ways in which data diversity should be used and how it should be integrated with design diversity are open questions.


ACKNOWLEDGMENT

It is a pleasure to acknowledge E. Migneault for thoughts about the failure regions and data diversity as a general concept. S. Brilliant's work in identifying the program faults and matching those faults with failure cases made it possible to carry out the empirical parts of this research. We thank D. Miller for providing a valuable critique of an earlier version of this work. We are also pleased to acknowledge L. Yount for a discussion about time-skewed inputs.

REFERENCES


- [1] P. Ammann, "Data diversity: An approach to software fault tolerance," Ph.D. dissertation, Univ. Virginia, Jan. 1988.
- [2] P. Ammann and J. C. Knight, "Data diversity: An approach to software fault tolerance," in *Dig. FTCS-17, Seventeenth Int. Symp. Fault Tolerant Comput.*, Pittsburgh, PA, July 1987, pp. 122-126.

- [3] A. Avizienis, "The *N*-version approach to fault-tolerant software," *IEEE Trans. Software Eng.*, vol. SE-11, Dec. 1985.
- [4] S. S. Brilliant, J. C. Knight, and N. G. Leveson, "Analysis of faults in an *N*-version software experiment," Univ. Virginia Tech. Rep. TR-86-20, Sept. 1986.
- [5] F. Cristian, "Exception handling," in *Resilient Computing Systems*, Vol. 2, T. Anderson, Ed. New York: Wiley, to be published.
- [6] J. Gray, "Why do computers stop and what can be done about it?" Tandem Tech. Rep. 85.7, June 1985.
- [7] J. C. Knight and N. G. Leveson, "A large scale experiment in *N*-version programming," *IEEE Trans. Software Eng.*, vol. SE-12, Jan. 1986.
- [8] D. J. Martin, "Dissimilar software in high integrity applications in flight control," in *1982 AGARD Conf. Proc.* 330, Software for Avionics, pp. 36-1-36-13.
- [9] M. A. Morris, "An approach to the design of fault tolerant software," M.Sc. thesis, Cranfield Instit. Technol., Sept. 1981.
- [10] P. M. Nagel and J. A. Skrivan, "Software reliability: Repetitive run experimentation and modeling," NASA Rep. CR-165836, Langley Res. Cent., Feb. 1982.
- [11] B. Randell, "System structure for software fault tolerance," *IEEE Trans. Software Eng.*, vol. SE-1, June 1975.
- [12] *Tutorial: Software Testing and Validation Techniques*, 2nd ed. IEEE Comput. Soc. Press, 1981, pp. 347-348.



Paul E. Ammann received the A.B. degree in liberal arts from Dartmouth College, Hanover, NH, in 1983, and the M.S. and Ph.D. degrees in computer science from the University of Virginia, Charlottesville, in 1985 and 1987, respectively.

Currently, he is employed by the Software Productivity Consortium, Reston, VA. His current interests include software reliability and software fault tolerance.



John C. Knight received the B.Sc. degree in mathematics from the Imperial College of Science and Technology, London, England, and the Ph.D. degree in computer science from the University of Newcastle-upon-Tyne, Newcastle-upon-Tyne, England in 1969 and 1973, respectively.

From 1974 to 1981 he was with NASA's Langley Research Center and in 1981 he joined the University of Virginia as an Associate Professor of Computer Science. At present he is with the Software Productivity Consortium, Reston, VA.

His research interests include software reliability and software fault tolerance.